

INTERNAL

AI System Description Template

Template - issued 15.05.26

Status	Draft
Security	Company Restricted
Version No.	1.0
Approved by	Accountable Owner / Approver
Author(s)	AI Governance Team

Contents

- 1. How to use this template 3
 - 1.1 System Identity.....3
 - 1.2 Summary3
- 2. System Overview (C4) 4
 - 2.1 Business Context and Objectives4
- 3. Intended Use and Users (C9)..... 5
 - 3.1 Problem Statement.....5
 - 3.2 Rationale for Using AI5
 - 3.3 Success Criteria5
- 4. Architecture and Components (C4) 6
 - 4.1 Component Inventory.....6
 - 4.2 Tooling and Approval Status6
- 5. Model and Component Selection Rationale (C4) 8
- 6. Data Dependencies 9
- 7. Logging and Telemetry (C7)..... 10
- 8. User and Operator Information (C9) 11
- 9. Known Limitations and Failure Modes (C9)..... 12
- 10. Non-Functional Requirements..... 13
- 11. Retirement Plan 14
- 12. Maintenance and Support..... 15
- 13. Change History 16

1. How to use this template

This template captures the AI System Description (Model/System Card) for an AI system delivered or operated by Kainos. It is one of the four mandatory artefacts defined in the Kainos AI Framework Policy and is the single source of truth for what the system is, how it works, what it uses, and its limitations. It supports controls C4, C7, and C9.

Complete every section. Use [square-bracketed grey text] as fillable placeholders. Cross-reference the Risk and Impact Assessment, Data Management Record, and Monitoring and Review Record for the same system. *Italic guidance text in this template should be removed once the section is completed.*

1.1 System Identity

Field	Value
AI System Name	[Enter system name]
OneTrust Assessment	[OneTrust Assessment ID]
Risk Tier	[Tier 2 Medium / Tier 3 High]
Customer (if applicable)	[Enter customer name]
Data Governance Role	[Controller / Joint Controller / Processor]
Data Controller	[If not Kainos]
Monitoring Owner	[Customer / Kainos / Shared]
Accountable Owner	[Name, role]
Approver	[Name, role - for Tier 2/3 not the Owner]
Reviewer(s)	[Security / Data Ethics / QA / Legal as required]
Lifecycle Stage	[Discover / Design / Develop / Deploy & Operate]
Version	[x.y]
Date	[dd.mm.yy]
Status	[Draft / Approved / Superseded]

1.2 Summary

Key points to capture in this artefact:

- What the system does and who it is for.
- How it is built (architecture, components, models).
- Why those choices were made (selection rationale).
- What it depends on (data, services, sub-processors).
- What it logs and how that logging supports oversight.
- What users and operators must know to use it safely.
- Known limitations and out-of-scope uses.

2. System Overview (C4)

Provide a concise plain-English description of the system: the business problem it solves, the value it delivers, and the boundary of the system.

[Describe the AI system at a level a non-technical reader can follow.]

2.1 Business Context and Objectives

State the business objectives the system supports, expected benefits, and success measures.

[Enter business context and measurable objectives.]

3. Intended Use and Users (C9)

Describe the intended use cases, the users and operators, and any explicit out-of-scope uses.

Aspect	Description
Primary intended use	<i>[What the system is designed to do.]</i>
Primary users	See Risk & Impact Assessment section 2 for Stakeholder and Impact Analysis
Operators / oversight role	<i>[Who runs, monitors, and is accountable.]</i>
Out-of-scope uses	<i>[Uses the system must not be applied to.]</i>
Deployment context	<i>[Internal / customer-facing; region; channel.]</i>

3.1 Problem Statement

[State the problem the system is intended to solve.]

3.2 Rationale for Using AI

[Explain why an AI approach is appropriate; reference alternatives considered.]

3.3 Success Criteria

[Measurable criteria for success. This should cover business outcomes only, e.g. Reduce processing time by 20%, improve customer satisfaction. Technical acceptance criteria is covered in Artifact 4 section 3 – Pre-stated Acceptance Criteria.]

4. Architecture and Components (C4)

What to include:

All the AI components themselves (model(s), orchestrator, vector store, guardrails, agent tools).

Any non-AI component that touches the AI - i.e. feeds its inputs, consumes its outputs, or sits inside its trust boundary. Examples: the data source/RAG pipeline feeding it, the auth layer in front of it, the downstream service that acts on its decisions.

What you can leave out:

Non-AI parts of the solution that don't interact with the AI (e.g. an unrelated billing or reporting module). For those, reference your normal solution-architecture documentation rather than duplicating it. You do not need to inventory every component of the entire solution.

[Insert architecture diagram or reference.]

4.1 Component Inventory

Component	Type	Provider / Version	Hosting / Region	Role in the system
[e.g. Foundation model]	[LLM / ML / Rule-based]	[Vendor, model name, version]	[Region]	[Function in flow]
[e.g. Vector store]	[Database / Index]	[Vendor, version]	[Region]	[Function in flow]
[e.g. Orchestrator]	[Framework]	[Library, version]	[Region]	[Function in flow]
[Add rows as needed]				

4.2 Tooling and Approval Status

List all tooling used to develop, evaluate, monitor, and govern the AI system. For each tool, record its status against the AI Governance approved tooling list. Where a tool is not on the approved tooling list, reference the relevant Approvals & Exceptions Log entry authorising its use. Where the use of a tool is mandated or approved by the customer, record the status as "Customer Approved" and include any supporting customer approval or contractual requirement where available.

Tool / platform	Category	Approved-list status	Approval reference
[e.g. Azure OpenAI]	Model gateway / inference	[Approved / Pending / Exception / Customer Approved]	[Approved-list ID or Approvals log ref]
[e.g. MLflow]	Experiment tracking / model registry	[Approved / Pending / Exception / Customer Approved]	[Ref]
[e.g. promptfoo / DeepEval]	Evaluation framework	[Approved / Pending / Exception / Customer Approved]	[Ref]
[e.g. Guardrails AI / NeMo Guardrails]	Guardrails / safety library	[Approved / Pending / Exception / Customer Approved]	[Ref]

Tool / platform	Category	Approved-list status	Approval reference
[e.g. Azure Monitor / Datadog]	Monitoring / observability	[Approved / Pending / Exception / Customer Approved]	[Ref]
[Add rows as needed]			

5. Model and Component Selection Rationale (C4)

Record the rationale for each material model, tool, or component used by the system. Where a component was selected by Kainos, document alternatives considered and the rationale for selection. Where a component was mandated by the customer, regulator, or existing enterprise architecture, record the source of the requirement.

Decision	Selection Owner	Options considered	Selected option	Rationale (incl. trade-offs)
[Foundation model choice]	[Kainos / Customer / Joint]	[Options]	[Selected]	[Why; what was traded off]
[Hosting model]	[Kainos / Customer / Joint]	[Options]	[Selected]	[Why]
[Guardrail approach]	[Kainos / Customer / Joint]	[Options]	[Selected]	[Why]
[Add rows as needed]				

6. Data Dependencies

List the data the system relies on (training, fine-tuning, retrieval, evaluation, runtime input/output). Full data details belong in the Data Management Record - reference it here.

Use	Data set / source	Reference in DMR
[Training / Fine-tuning]	[Source]	[DMR Section #]
[Retrieval (RAG)]	[Source]	[DMR Section #]
[Runtime input]	[Source]	[DMR Section #]
[Evaluation / test sets]	[Source]	[DMR Section #]

7. Logging and Telemetry (C7)

Describe how inputs, outputs, tool/agent actions, and material decisions are logged so that events can be reconstructed for the retention period defined in the Data Management Record. Cover what is captured, where it is stored, how access is controlled, and the retention period.

Aspect	Description
What is logged	[Prompts, completions, tool calls, decisions, model version, user ID, session ID, latency, errors.]
Storage location	[Service / region / tenancy.]
Retention period	[N days / months; reference DMR rule.]
Access controls	[Roles authorised, audit trail.]
PII / sensitive handling	[Redaction, encryption, lawful basis.]
Reconstruction capability	[How to replay an event from logs for audit or incident.]

8. User and Operator Information (C9)

Capture the information that must be made available to users and operators. This may be published in product help, an in-product notice, or operator runbooks - record where in the row that applies.

Information	Content	Where published
Intended use and audience	[Plain-English statement of intended use.]	[Help page / in-product]
Limitations	[What the system cannot reliably do and how are any limitations communicated to users.]	[Help page]
Required oversight	[Where human review is required.]	[Operator runbook]
Redress channel	[How to raise issues, request review, opt out.]	[Support channel]
Data handling notice	[Lawful basis, retention, data subject rights.]	[Privacy notice]
AI Disclosure	[How are users informed that AI is being used]	[In product / help page]

9. Known Limitations and Failure Modes (C9)

List the known limitations and failure modes of the system. This is essential context for users, operators, and reviewers.

Limitation / failure mode	Why it occurs	Mitigation in place
[e.g. Hallucination on niche topics]	[Cause]	[RAG, citations, guardrails]
[e.g. Stale knowledge]	[Training cut-off]	[Refresh cadence]
[e.g. Bias in outputs]	[Cause]	[Bias testing, retraining]
[Add rows as needed]		

10. Non-Functional Requirements

Non-functional requirements fall into two groups. **Standard IT NFRs** already governed by normal project delivery (performance, availability, scalability, baseline security, accessibility), should be **referenced** from the project's requirements/NFR artefact rather than re-typed here.

This section captures the **AI-specific** quality requirements: the ones that link to the Risk & Impact Assessment (C5) and are tested under C6. Record each with a threshold, an owner, and how it's verified. Tier 2 and Tier 3 list explicit thresholds with owners.

The requirements should reflect the system's context and risk. You can legitimately mark an NFR **N/A with justification** (e.g. a deterministic rule-based component = "drift not applicable"; no protected-group impact = "fairness NFR not applicable").

NFR area	Requirement / threshold	Owner	Verification
Model performance / accuracy	[e.g. ≥ 0.85 F1 on holdout; task-specific metric]	[Owner]	[Eval report ref - C6]
Robustness	[Performance holds within X% under adversarial inputs / foreseeable misuse]	[Owner]	[Robustness / red-team test - C6/C14]
Fairness / bias	[Max disparity across protected groups $\leq X$]	[Owner]	[Bias evaluation ref - C6]
Explainability	[Required explanation type and audience, e.g. feature attributions for reviewers]	[Owner]	[Review sign-off - C9]
Drift tolerance	[Retrain/alert trigger, e.g. $>X\%$ metric drop over period]	[Owner]	[Monitoring ref - C10]
AI-specific security	[Prompt-injection and data-exfiltration controls]	[Owner]	[Control evidence - C7]
Standard IT NFRs	Performance, availability, scalability, accessibility (WCAG 2.2 AA where user-facing) - see [project NFR/requirements doc]	[Owner]	Per normal delivery

11. Retirement Plan

Document how the system will be retired. Cover the trigger conditions, data and model disposition, stakeholder communications, and where retirement evidence will be retained. Tier 2 and Tier 3 must record each row below explicitly.

Aspect	Detail
Retirement triggers	[e.g. superseded model, supplier exit, regulatory change, low usage]
Data disposition	[Delete / archive / migrate; retention reference in Data Management Record]
Model and artefact archive	[Where the final model, prompts, weights, and evaluation evidence will be retained]
Stakeholder communications	[Users, operators, customers, regulators (if applicable); timing]
Successor / fallback	[Replacement system or manual process; cutover plan]
Evidence retention	[Where retirement record is stored; retention period]

12. Maintenance and Support

Record who is responsible for routine maintenance, model refreshes, and operational support, including out-of-hours arrangements where applicable.

Aspect	Description
Maintaining team	[Team and contact]
Support model	[Hours, SLAs]
Refresh / retraining cadence	[Cadence and trigger]
Dependencies on suppliers	[Suppliers and contract refs]

13. Change History

Record material updates to this AI System Description. Minor edits do not require a row.

Version	Date	Author	Change
[1.0]	[dd.mm.yy]	[Name]	[Initial issue.]
[]	[]	[]	[]

kainos®

