

Review of Model-Provider Terms Foundation Models on Amazon Bedrock

1 Background to the Request

The Probabilistic Triage team are exploring how probabilistic AI could support NHS triage alongside existing deterministic triage capability. The team is undertaking controlled discovery/Alpha work and wishes to evaluate foundation models obtainable through Amazon Bedrock, specifically:

- Anthropic Claude Opus 5,
- OpenAI GPT-5.3 (or the applicable OpenAI model/version available through Bedrock), and
- Amazon Nova Pro.

The proposed *initial* use is in customer development and test environments only.

The models would be evaluated as enabling tools for activities such as generating and refining candidate Bayesian-network structures and supporting artefacts, producing synthetic clinical test vignettes/test cases, assisting with test-harness development, and supporting analysis of model behaviour.

Any clinically meaningful artefact produced with model assistance would remain subject to human review, clinical assurance and explicit clinician approval gates before it could influence a released probabilistic triage model. The intention is not, at this stage, to allow a general-purpose foundation model to provide autonomous patient-facing triage advice or make an unreviewed clinical decision.

Personal data:

- Initial exploration should be capable of being conducted using synthetic or appropriately de-identified/non-live data.
- No production patient data is required for the *initial* model comparison.
- Any future proposal to use identifiable or production clinical data would be subject to separate information-governance, security, clinical-safety and legal assessment.

2 Legal Request

Legal advice was requested to establish whether NHS England can approve controlled evaluation and potential future use of the named foundation models through Amazon Bedrock for Probabilistic Triage, and to identify any conditions that must be met before model access is enabled.

The scope of the review was limited to the product specific terms and conditions which are available online¹.

3 Legal Team Approach

The request received was broad and did not include either the commercial agreement that had been entered into, nor the specific contract areas by product. The request was to focus on the product specific terms. The legal team have reviewed the documentation available online, but since it has not been provided with the relevant NHS England AWS Bedrock core documentation certain assumptions have been made. We recommend that the contract documentation is clarified before reliance placed on advice of an incomplete contract set. The AWS contract layer is relevant for the review and should be considered and reviewed in addition to what is set out here.

Annex 1 to this report includes the links associated with the models. The legal team are not able to complete the contract journey and therefore request that these are also reviewed and confirmed by the Probabilistic Triage Team.

Products in scope for phase 1 are the following foundation model provider's terms:

- Anthropic Claude Opus 5;

Products in scope for phase 2 are the following foundation model provider's terms:

- OpenAI GPT-5.6 Sol (substituted for GPT-5.3, which we could not locate on AWS Bedrock);
- Amazon Nova Pro (Nova 2 Pro is preview-only and not reviewable) and
- Amazon Nova 2 Lite

The initial phase uses synthetic/de-identified data, but the review assesses each product as if personal and special-category data were present, so findings remain valid for later phases. Where a finding depends on that assumption, it is stated.

The review drew a distinction between trial & free-service terms, and fully paid standard service terms.

The legal review sought to understand the terms and conditions in the following topic areas, as follows:

Purpose / IP:

- IPR – customer content (inputs)
- IPR – anonymised / de-identified data; provider telemetry
- IPR – outputs
- Purposes for which the foundation model provider may use customer inputs and customer outputs; and any restrictions.
- Whether there can be any training or product improvement based on customer data.
- Scope of the foundation model provider's access to prompts and outputs on AWS Bedrock.
- Foundation model provider's responsibility for IPR infringement in outputs

Use Case Constraints

- Provisions constraining NHSE's acceptable use – healthcare / clinical (vs stated use)

¹ Website links reviewed: 4 September 2026

Confidentiality and security:

- Any commitments to known security certifications and standards
- Confidentiality of inputs and outputs
- Any customer audit rights

Commercial and liability

- Accuracy warranties / disclaimers
- Charges not fully detailed; unexpected bills and charges
- Foundation model provider's liability cap
- Obligations on the customer to provide indemnities to foundation model provider.
- Scope of any usage reporting
- Foundation model provider's responsibility for IPR infringement in outputs

Data Protection:

- The controller / processor status of the foundation model provider.
- Controls relating to sub-processors of the foundation model provider (if any).
- Any information relating to sovereignty / residency / international transfers
- Whether a human review of content (if any) can be switched off.
- Whether there is any retention of inputs/outputs by the foundation model provider.
- Any information around the deletion of customer data

Duration, interruption to service risk:

- Model deprecation and lifecycle
- Foundation model provider suspension rights
- Foundation model provider termination rights
- Customer termination rights

Other Issues

- Any other issues of relevance identified during review not covered by the above areas: e.g. publicity clauses.

4 High Level Review: Phase 1

The review has focussed initially on Anthropic Claude Opus 5.

On AWS Bedrock, AWS runs the model and for all three providers states that retained content is never shared with the model provider, so the implication being that 'data leakage to the provider' should be largely resolved at the platform level. Since this is not a conclusion solely from the contractual terms and conditions this should be verified by the Probabilistic Triage Team and our technical team. Our legal conclusion in some areas is subject to this fact verification.

The processing location is a per-model AWS Bedrock fact, not a contractual promise, and it differs between the four models. NHS England would need to do a review of the processing jurisdictions before permitting such processing of any personal data. If relevant, Anthropic operates an EEA processing. The product specific contract terms state that Claude Opus 5 on AWS Bedrock operates on a zero-data-retention basis by default. On termination, Anthropic's DPA appears to require deletion of Customer Data within 30 days, although Anthropic may retain information where necessary to comply with legal obligations, resolve disputes, or address harmful-use concerns. This would also permit retention where there is a concern regarding legal disputes which is a low bar. Separately, the contract requires Confidential Information to be returned or destroyed on request. Presumably such caveats are irrelevant though if the platform model means that Anthropic do not receive the data.

The Anthropic IP indemnity applies to "paid use" only and not to the trial licence (clause K.1). Otherwise, everything else is identical for trial and paid use. The Customer retains ownership of inputs and Anthropic disclaims any ownership claim for outputs (which may be the maximum it can state under some IP regimes).

All the indemnification made by the customer are uncapped, which includes any third-party claims relating to (a) Inputs or other data provided by Customer, or (b) use of the Services in violation of the Usage Policy, the Service Specific Terms, or Section D.4 (Use Restrictions). An uncapped indemnity to a supplier is a contingent liability that may need HM Treasury / Managing Public Money consideration. This is a contingent liability that may need HM Treasury Managing Public Money consideration. This should be flagged to Finance for consideration.

5 Recommendations

There are five legal recommendations:

- Review our linked documents and confirm these are the appropriate set. Commission a full review of terms including the AWS Bedrock terms. Currently we have only reviewed the product specific terms and conditions, and the AWS Bedrock terms may take higher priority and vary the position or introduce new issues.
- Team to consider the outlined use case and to confirm it does not breach the healthcare High-Risk requirements. We do not think it does as the stated restrictions are framed by reference to individuals/consumers and the use case here is internal tooling with clinician gates, so the requirements appear to be met, but please review to ensure we have understood correctly. (see row 8 in Annex 2 of this report).
- Team to consider the outlined use case and to confirm it does not breach the prohibition on using outputs to train an AI model without authorisation. Could model-assisted refinement of a Bayesian-network structure be argued to fall within it? (see row 8 in Annex 2 of this report).
- Work with PTT to develop an AWS Bedrock all-models system wide DPIA in good time. Review the contractual data protection sections in more detail prior to using any personal data on the system. On the basis of our brief review, particularly given the capability to switch between models, there appears to be enough substance to warrant a fuller review to consider personal data and special category data use cases.
- Raise the uncapped customer indemnity issue with Finance to see if this needs approval in line with HM Treasury Managing Public Money guidance.

Author; Hazel Randall

Title: Solicitor and Deputy Director of Legal Services, Commercial Information and Technology Domains.

Version: 1.0 Draft for comment.

Dated 7th September 2026

Annex 1 – T&Cs Website Links Identified

Anthropic – Claude Opus 5

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
Anthropic on Bedrock – Commercial Terms of Service	The contract with Anthropic. AWS states Anthropic services on Bedrock are "sold by Anthropic" and use means agreement to these terms. UK customer contracts with Anthropic Ireland, Limited; Irish law; arbitration in Dublin (cl. M.7, J.2). Incorporates the four documents below by reference (cl. D.2, M.9).	Version: August 18, 2025	https://aws.amazon.com/legal/bedrock/third-party-models/ (Anthropic section)	Located
Anthropic Usage Policy (AUP)	Acceptable use. Incorporated by cl. D.2(a). Contains Universal Usage Standards, High-Risk Use Case Requirements (stated to apply to consumer-facing use cases) and Additional	Current version effective September 15, 2025 (per Anthropic announcement); confirm on page	https://www.anthropic.com/aup (terms link to https://console.anthropic.com/legal/aup)	Located

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
	Use Case Guidelines. Key document for review row 8.			
Anthropic Data Processing Addendum	Incorporated by cl. C. Governs personal data processing; expected to contain UK GDPR processor terms, sub-processor list reference and transfer mechanism. Key for rows 1–7, 25, 26.	Not retrieved	https://www.anthropic.com/legal/data-processing-addendum	Located
Anthropic Service Specific Terms	Incorporated by cl. D.2(c). May contain Bedrock-specific or feature-specific terms.	Not retrieved	https://www.anthropic.com/legal/service-specific-terms	Located
Anthropic Supported Regions Policy	Incorporated by cl. D.2(b). Confirms UK is a supported customer location.	Not retrieved	https://www.anthropic.com/supported-countries	Located
Claude Opus 5 – Bedrock model card (AWS)	Technical facts the review depends on: launch 24 Jul 2026, lifecycle Active, EOL N/A, regional availability and inference-profile options, zero-data-retention default,	Page current as of retrieval	https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-anthropic-claude-opus-5.html	Located

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
	billing via AWS Marketplace under the model provider.			
Claude in Amazon Bedrock – Anthropic platform docs	Anthropic's own description of the Bedrock offering; states data handling for this offering is governed by Amazon Bedrock and points to AWS data-protection docs.	Page dated ~June 2026	https://platform.claude.com/docs/en/build-with-claude/claude-in-amazon-bedrock	Located
Claude Opus 5 system card (transparency material)	Anthropic publishes a system card per model. Not a contract; transparency material for row 15 and clinical-safety colleagues.	Not retrieved	Locate on anthropic.com – URL not verified	To locate
DPIA	None published by Anthropic. Expected: the NHS body is controller; the DPA is the relevant artefact.	n/a	n/a	None found

OpenAI - GPT-5.6 Sol (or Terra)

GPT-5.3 is not on Bedrock. See Flag 1. The table below is prepared on the assumption that the substitute is GPT-5.6 Sol (or Terra); the same terms set applies to GPT-5.5 and GPT-5.4.

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
OpenAI section of "Serverless Third-Party Models on Amazon Bedrock"	The page every OpenAI model card links to for "End User License Agreements and Terms of Use". This section determines whether OpenAI models are "sold by OpenAI" under an OpenAI EULA, or "sold by AWS" under additional terms. I could not retrieve this section (page too long for automated fetch) – it must be read manually.	Page last updated 24 Aug 2026 (page metadata)	https://aws.amazon.com/legal/bedrock/third-party-models/#amsc13--xttmgl	NOT RETRIEVED – manual read required
OpenAI Business Terms and Usage Policies	Candidate documents only. Include in the set ONLY if the section above incorporates them. If OpenAI models are "sold by AWS", OpenAI's own terms may not apply at all.	Business Terms / Service Terms pages dated June 2026 on openai.com	https://openai.com/policies/business-terms/ ; https://openai.com/policies/usage-policies/ ; https://openai.com/policies/service-terms/	Conditional – applicability unconfirmed

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
GPT-5.6 Sol – Bedrock model card (AWS)	Technical facts: launch 13 Jul 2026, Active, EOL N/A; regional availability (see Flags); on bedrock-runtime from London only Global cross-region inference is available; pricing table shows In-Region / Geo / Global tiers.	Page current as of retrieval	https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-openai-gpt-56-sol.html	Located
OpenAI models in Amazon Bedrock – AWS index of OpenAI model cards	Lists the OpenAI models actually on Bedrock: GPT-5.6 Sol/Terra/Luna, GPT-5.5, GPT-5.4, two Daybreak cyber models (restricted), and four open-weight gpt-oss models. No GPT-5.3.	Page current as of retrieval	https://docs.aws.amazon.com/bedrock/latest/userguide/model-cards-openai.html	Located
OpenAI developer guide – "OpenAI models in Amazon Bedrock"	OpenAI's own statement of data handling on Bedrock: AWS does not share request/response content with OpenAI in default or none retention mode; default mode retains classifier-flagged traffic up to 30 days and Responses API store:true responses for 30 days; full	Feature table dated 13 Jul 2026	https://developers.openai.com/api/docs/guides/amazon-bedrock	Located

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
	ZDR requires AWS approval per account and model; AWS Regions differ from OpenAI data-residency jurisdictions.			
GPT-5.6 system card (transparency material)	Linked from the AWS model card as the model/service card.	Not retrieved	https://deploymentsafety.openai.com/gpt-5-6/gpt-5-6.pdf	Located
DPIA	None found.	n/a	n/a	None found

Amazon – Nova Pro

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
AWS Service Terms – relevant sections	Nova is AWS first-party, so its terms are the AWS Service Terms (reviewed as the agreed exception). Sections identified: 1.23 (AI/ML services must comply with the AWS Responsible AI Policy); 50 (Amazon Bedrock, including 50.3 on non-use of inputs/outputs for training per AWS Nova docs, 50.10 uncapped IP indemnity for Indemnified Generative AI Services including Amazon models, 50.12 Third-Party Content). Section numbers to be verified against the live page.	Page updated within days of retrieval (search metadata "3 days ago")	https://aws.amazon.com/service-terms/	Located; section 50 not retrieved in full
AWS Responsible AI Policy	Supplements the AUP and Service Terms for all AI/ML services. Prohibitions plus a	Last Updated: January 13, 2025	https://aws.amazon.com/machine-learning/responsible-ai/policy	Located

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
	"consequential decisions" requirement (human oversight, testing, safeguards) expressly naming medical diagnosis and medical advice. Key for row 8.			
Nova Pro – Bedrock model card (AWS)	Technical facts: launch 5 Dec 2024, lifecycle Active but EOL "no sooner than 12/4/2025" (date already passed); In-Region availability in eu-west-2 (London); EU geo profile routes across EU regions only; abuse detection listed as not supported for this model.	Page current as of retrieval	https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-amazon-nova-pro.html	Located
AWS AI Service Card – Amazon Nova Micro, Lite, Pro and Premier	Transparency material; states AWS does not use Bedrock inputs/outputs to train Nova; abuse detection is automated with no human review; customers in high-risk workflows must implement human oversight.	Not retrieved in full	https://docs.aws.amazon.com/ai/responsible-ai/nova-micro-lite-pro/overview.html	Located

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
Amazon Nova User Guide – Responsible use	States Bedrock does not store or review prompts/completions and does not use them to train Nova; cites Service Terms s.50.3.	Not retrieved in full	https://docs.aws.amazon.com/nova/latest/userguide/responsible-use.html	Located
Amazon Bedrock model lifecycle page	AWS commitment that a model stays on Bedrock at least 12 months after launch and Legacy status lasts at least six months before EOL. Relevant to row 16 for all three products.	Page current as of retrieval	https://docs.aws.amazon.com/bedrock/latest/userguide/model-lifecycle.html	Located
DPIA	None found.	n/a	n/a	None found

AWS Bedrock platform documents

These are AWS-layer documents and outside the contractual scope, but each product's data-handling position points to them, so they are listed here so the Phase 2 reviewer can cite them as the source of a fact rather than guessing.

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
Amazon Bedrock – Data protection	AWS-layer, but the Anthropic docs expressly defer data handling on Bedrock to this page, so the Anthropic review cannot be completed without it.	Not retrieved	https://docs.aws.amazon.com/bedrock/latest/userguide/data-protection.html	Located
Amazon Bedrock – Data retention	Defines retention modes (default / none) and ZDR eligibility. Needed for rows 3–5 on all three products.	Not retrieved	https://docs.aws.amazon.com/bedrock/latest/userguide/data-retention.html	Located
Amazon Bedrock – Abuse detection	Per-model list of what is retained for abuse detection and for how long. Needed for row 4.	Not retrieved	https://docs.aws.amazon.com/bedrock/latest/userguide/abuse-detection.html	Located
Amazon Bedrock – Cross-Region inference / regional	Source-to-destination region tables for each inference profile. Needed to answer row 6 precisely.	Not retrieved	https://docs.aws.amazon.com/bedrock/latest/userguide/models-region-compatibility.html	Located

Document	What it is / why it is in the set	Version / date	URL	Status (4 Sep 2026)
availability by model				

Annex 2 – Phase 1

Anthropic – Claude Opus 5

Contract: Anthropic on Bedrock – Commercial Terms of Service, version 18 August 2025, incorporating the Usage Policy (effective date 15 Sep 2025), DPA (effective date 24 Feb 2025 – version to confirm), Service Specific Terms and Supported Regions Policy.

Document	Version / date	URL
Anthropic on AWS Bedrock – Commercial Terms of Service	August 18, 2025	https://aws.amazon.com/legal/bedrock/third-party-models/ (Anthropic section)
Anthropic Usage Policy	Effective September 15, 2025	https://www.anthropic.com/legal/aup
Anthropic Data Processing Addendum	Effective 24 Feb 2025 (per secondary source – confirm)	https://www.anthropic.com/legal/data-processing-addendum
Anthropic Service Specific Terms	Not confirmed	https://www.anthropic.com/legal/service-specific-terms
Anthropic Supported Regions Policy	Not confirmed	https://www.anthropic.com/supported-countries
Anthropic on AWS Bedrock Commercial Terms – Dec 2023 (superseded; cited for usage-data wording)	Dec 2023	https://www-cdn.anthropic.com/6b68a6508f0210c5fe08f0199caa05c4ee6fb4dc/Anthropic-on-Bedrock-Commercial-Terms-of-Service_Dec_2023.pdf

Document	Version / date	URL
Claude Opus 5 – Bedrock model card	Launch 24 Jul 2026	https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-anthropic-claude-opus-5.html
Claude in Amazon Bedrock – Anthropic docs	~June 2026	https://platform.claude.com/docs/en/build-with-claude/claude-in-amazon-bedrock
AWS What's New – Claude Opus 5 on AWS	24 Jul 2026	https://aws.amazon.com/about-aws/whats-new/2026/07/claude-opus-5-aws/

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
1	IPR – customer content (inputs)	Customer retains all rights to Inputs; Anthropic disclaims rights received under the Terms (AT cl. B).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Green.	
2	IPR – anonymised / de-identified data; provider telemetry	The Customer's own de-identified inputs are Inputs and remain the Customer's (AT cl. B). The 2025 AT does not define or allocate usage data/telemetry (the 2023 version did – "Usage Data excludes Customer Content").	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Gap: the 2025 terms drop the express carve-out that usage data excludes Customer Content.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		Feedback may be used by Anthropic without obligation (AT cl. A.3).		Ask for confirmation. Feedback clause is standard – we should avoid sending any personal data including clinical content as "feedback".	
3	IPR – outputs	Customer owns Outputs; Anthropic assigns any right, title and interest in Outputs to Customer subject to compliance (AT cl. B).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Green. Assignment is conditional on compliance – another reason to settle row 8.	
4	Purposes for which provider may use inputs/outputs; restrictions	Anthropic disclaims any rights it receives to Customer Content 'Inputs' and 'Outputs' and may use it only to exercise rights/perform obligations as Confidential Information (Anthropic on Bedrock Commercial Terms v. 18 Aug 2025 ("AT"), clause B, E.2 on confidentiality). It specifically states "Anthropic may not train models on Customer Content from Services. " "Customer Data" means all data or other information submitted through the	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Strong. There is a residual question what usage/metadata Anthropic receives from AWS, which the 2025 terms no longer describe. Ask AWS/Anthropic to confirm in writing.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		Services by or for Customer or a Customer Affiliate. Stated that “ the technology provided by Anthropic to enable access to the Hosted Services does not give Anthropic access to any Customer instance within the applicable Cloud Platforms, including the Inputs or Outputs they contain” (clause D2) This is subject to permissions granted in subject to complying with DPA provisions, see Part B of Schedule 1 for scope. Personal data processed under the Anthropic DPA (cl. C). Anthropic's own docs state data handling for the Bedrock offering is governed by Amazon Bedrock (Claude in Amazon Bedrock docs). AWS states retained abuse-detection data is not shared with third-party model providers (Bedrock abuse-detection page).			
5	Training / product improvement on customer data	Express prohibition: Anthropic may not train models on Customer Content from the Services (AT cl. B).	No distinction, except IP indemnity applies only	Green. Clear contractual prohibition, reinforced by AWS.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		AWS separately states neither AWS nor third-party model providers use Bedrock inputs/outputs to train models (Bedrock FAQ).	to "paid use" (Anthropic Terms cl. K.1).		
6	Provider access to prompts and outputs on Bedrock	Bedrock operates zero operator access and, for Opus 5, zero data retention by default (AWS Opus 5 launch notice 24 Jul 2026; abuse-detection page – Opus 5 is not among the models for which traffic is retained). AWS states retained data is never shared with model providers. Net effect: we believe this means Anthropic does not receive prompts or outputs, but would appreciate confirmation that the definition does not effectively mean ‘a subset of customer provided or processed data is not shared’. Note: the Dec 2023 version of these terms expressly said Anthropic receives only usage data from AWS excluding	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Green on the stated facts, but the assurance here comes from AWS documentation rather than the Anthropic contract. We would like to see a contractual commitment	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		Customer Content; the 18 Aug 2025 version is silent on usage data, so has this changed?			
7	Provider responsibility for IPR infringement in outputs	Anthropic defends and indemnifies for third-party IP claims arising from paid use in accordance with the Terms, including Outputs (AT cl. K.1). Exclusions: modifications, combinations, Customer Inputs, knowing infringement, patented inventions in Outputs, trademark use (cl. K.3). Indemnity is outside the liability cap (cl. L.3.b).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Good for the market; note it is limited to paid use and excludes Customer-Input-driven claims. Adequate.	
8	Acceptable use – healthcare / clinical (vs Stated Use)	AUP (effective date 15 Sep 2025) High-Risk Use Case Requirements cover healthcare decisions, diagnosis, patient care and medical guidance; they require a qualified professional to review content "directly affecting individuals or consumers" before finalisation, and AI	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Potentially red – two points. (1) The healthcare High-Risk requirements are framed around individuals/consumers; the Stated Use is internal tooling with clinician gates, so the requirements appear to be met rather than breached, but please	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		disclosure where outputs are presented directly to individuals. The intro states these apply to consumer-facing use cases. Universal standards also prohibit providing false or misleading medical/health information and misuse of health data without permission. Separately prohibited: using inputs and outputs to train an AI model without Anthropic's prior authorisation ("model scraping / distillation"), and building a competing product (AT cl. D.4).		review to ensure we have understood correctly. (2) The prohibition on using outputs to train an AI model without authorisation is a potential risk: question, could model-assisted refinement of a Bayesian-network structure could be argued to fall within it. If this is what we are doing we could seek Anthropic's written confirmation that the Stated Use is permitted before any output influences a developed triage model.	
9	Security certifications and standards	AT is silent on certifications. DPA Schedule 2 sets technical and organisational measures and refers to annual third-party audits such as SOC 2 (secondary source; verify). Bedrock platform certifications are AWS layer.	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Adequate for dev/test. For later phases obtain the SOC 2 report and DPA Schedule 2 through Anthropic; platform assurance sits with AWS.	
10	Confidentiality of inputs and outputs	Customer Content is Customer's Confidential Information; use limited to performing the Terms; disclosure only to	No distinction, except IP indemnity applies only	Standard and acceptable.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		need-to-know Representatives under equivalent obligations; reasonable care standard; legal-compulsion carve-out with notice; destruction on request save legal retention/backups (AT clause E.1–E.4).	to "paid use" (Anthropic Terms cl. K.1).	"Reasonable care" is the market norm; not the higher NHS standard, but according to statements on website the content never reaches Anthropic on Bedrock.	
11	Customer audit rights	AT is silent. DPA provides for audit reports and customer audits on conditions (secondary source; verify against DPA text).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Rely on the DPA. Verify the actual clause before citing.	
12	Accuracy warranties / disclaimers	Services and Outputs "as is"; all implied warranties disclaimed; no warranty of accuracy or completeness (AT cl. L.2). Customer must evaluate Outputs and notify its Users that factual assertions should be independently checked (cl. D.3).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms cl. K.1).	Standard for the sector. For clinical use the burden sits entirely with the NHS body's clinical-safety process (DCB0129/0160 – out of scope here).	
13	Charges not fully detailed; unexpected bills	AT says fees are at Anthropic's published rates and may change on 30 days' notice (clause H.1); in practice Opus 5 on Bedrock is billed through AWS Marketplace on the AWS bill under the model provider (model card, Pricing). Regional endpoints +10% (Claude docs).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Low risk in dev/test. Note the mismatch between AT clause H and AWS Marketplace billing should be reconciled in the AWS-layer review. This is needed to give the full legal picture.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		Taxes and gross-up on Customer (clause H.2).			
14	Provider liability cap	Each party's liability limited to fees paid in the previous 12 months; indirect and consequential loss excluded; applies in tort/negligence (AT clause L.3.a, c). Indemnities are excluded from the cap (clause L.3.b).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Cap is low (12 months' fees) and in dev/test that may be a trivial sum. Uncapped IP indemnity is the valuable part. Typical of hyperscaler-marketplace terms; unlikely to be negotiable via Bedrock. NHSE should note, but there is an element of recognise risk and manage.	
15	Customer indemnities to provider	Customer indemnifies Anthropic for third-party claims relating to (a) Customer Inputs or other data, or (b) use in breach of the AUP, Service Specific Terms or use restrictions (AT clause K.2). Uncapped (clause L.3.b).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Amber as a public body. All the indemnities made by the customer are uncapped, which includes any third party claims relating to (a) Inputs or other data provided by Customer, or (b) use of the Services in violation of the Usage Policy, the Service Specific Terms, or Section D.4 (Use Restrictions). An uncapped indemnity to a supplier is a contingent liability that may need HM Treasury / Managing Public Money consideration.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
				Flag to NHSE Finance.	
16	Usage reporting	AT is silent. Usage visibility is via AWS console/CloudWatch/CloudTrail (AWS layer; Claude on Bedrock docs recommend 30-day log retention).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Not addressed in product-specific terms – governed by AWS layer.	
18	Controller / processor status	The DPA is a controller-to-processor instrument (Module 2) with Anthropic as processor/importer (DPA clause I.1; Annex – secondary sources). On Bedrock, Anthropic does not receive Customer Content, so for content it is in practice neither controller nor processor; AWS is the processor under the AWS layer. The assumption that the provider is not a controller of customer content holds.	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Assumption confirmed for content. Anthropic may be a controller of the limited account/usage data it receives – immaterial for this review.	
19	Sub-processors	DPA defines Subprocessor (clause A.9) and provides for a sub-processor list with advance notice of changes (secondary sources; list reportedly at	No distinction, except IP indemnity applies only	Low materiality on Bedrock. Obtain the current sub-processor list for the file but note it will / may change.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		anthropic.com/subprocessors – not reviewed). On Bedrock the host is AWS under the customer's own AWS contract, not an Anthropic sub-processor.	to "paid use" (Anthropic Terms clause K.1).		
20	Sovereignty / residency / international transfers	Contracting entity for a UK customer is Anthropic Ireland, Limited; Irish law; arbitration in Dublin (AT preamble, clause M.7, J.2). Processing location is a Bedrock fact: Opus 5 has no In-Region option on bedrock-runtime anywhere; from eu-west-2 (London) the options are the EU geo profile (stated to keep data within EU regions – destination list not shown on the model card) or Global (worldwide). Regional endpoints carry a 10% premium (Claude on Bedrock docs). DPA incorporates EU SCCs Module 2/3 and the UK Addendum (DPA clause I.1, A.10 – secondary source).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Green. UK-only processing is not available. EU-region processing is a restricted transfer under UK GDPR, currently covered by the UK adequacy regulations for the EEA; document it in the ROPA. <u>Training needed internally</u> : Do NOT use the Global profile for anything beyond synthetic data. Verify the EU profile destination list and confirm London is a source region only. Could this be managed / turned off? Governing law Ireland; arbitration in Dublin under UNCITRAL rules (clause. J.2, M.7). Anthropic may amend the Terms on 30 days' notice, immediately for legal changes (clause M.3). These terms would need to be under constant review.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
21	Human review of content; can it be switched off?	Bedrock abuse detection is automated; human review by AWS is stated only for classifier-flagged Claude Fable 5/5.1 traffic and for suspected CSAM (abuse-detection page). Opus 5 is not in that list. Anthropic's AUP says its Safeguards team monitors for violations (AUP, intro) but on Bedrock it has no content to review. Not switchable by the customer; not applicable to Opus 5 on the facts.	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Not an issue for Opus 5 today. Note the contrast with the Fable models, where AWS retains all traffic 30 days with possible human review – do not assume the Opus position carries over if the model choice changes.	
22	Retention of inputs/outputs by provider	Zero data retention by default for Opus 5 (AWS launch notice; abuse-detection page). Under the DPA, on termination Anthropic deletes Customer Data within 30 days subject to exceptions for legal requirements, disputes and combating harmful use (DPA effective date 24 Feb 2025).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Green. For normal operation, customer prompts and generated outputs are not retained by Anthropic or AWS for future use, logging, model training or review. Zero data retention is a Bedrock configuration stated fact; confirm the effective retention mode on the actual account.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		Confidential <u>Information</u> destroyed on request, but excludes: “(i) Applicable Data Protection Laws or other applicable legal or regulatory requirements requires storage of the Customer Data, (ii) retention of the Customer Data by Anthropic is necessary to resolve a dispute between the parties, or (iii) retention of the Customer Data is necessary to combat harmful use of the Services.” (AT clause E.4, DPA clause H1b).		This therefore includes a limited product improvement associated with ‘<u>harmful use of the services</u>’. The "combat harmful use" deletion exception in the DPA is broad but presumably would be academic given zero data retention. Team to confirm. The default is delete within 30 days unless provider has a legitimate legal, regulatory or safety reason to keep some or all of it – which may include a concern regarding disputes.	
23	Deletion of customer data	Destruction of Confidential Information on request, except legal retention and automated backups which remain confidential (AT clause E.4). DPA: return/deletion within 30 days of termination with stated exceptions (DPA clause H.1.b).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Green on the facts if the website assertions are correct.	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
		On Bedrock with zero data retention there is nothing held to delete.			
25	Model deprecation and lifecycle	AT is silent. Bedrock policy: a model stays on Bedrock at least 12 months after launch; Legacy state lasts at least six months before EOL (Bedrock model-lifecycle page). Opus 5: launched 24 Jul 2026, Active, EOL N/A (model card).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	No contractual commitment beyond Bedrock's 12-month floor. For a triage model whose provenance may need to be reproduced, this may not be enough. Do we need to plan for replacement? Consider what is needed for evidencing methodology etc. The outputs must be verifiable on their own account, particularly if the model is deprecated and cannot reproduce results.	
25	Provider suspension rights	Anthropic may suspend if it reasonably believes there is a security risk, a breach of compliance/AUP/use restrictions, provision is unlawful or "would result in a material increase in the cost of providing the Services", or a vendor withdraws required third-party services; reasonable efforts to notify; no liability for suspension (AT clause I.3.a–b).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Amber. The "material increase in cost" ground is unusually wide and there is no liability for suspension. Could we mitigate architecturally rather than contractually?	

#	Review row	Finding – paid product (paraphrase, clause, source)	Trial / non-production	LEGAL REVIEWER POINTS	
26	Provider termination rights	Anthropic may terminate for convenience on 30 days' notice; for uncured material breach on 30 days' notice; immediately if provision is prohibited by law (AT clause I.2.a–c).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	30-day convenience termination is short for a regulated project but quite normal for API terms. Could we mitigate with model-agnostic architecture?	
27	Customer termination rights	Customer may terminate at any time for convenience with notice, or for uncured material breach on 30 days' notice (AT clause I.2.a–b).	No distinction, except IP indemnity applies only to "paid use" (Anthropic Terms clause K.1).	Green.	
28	Other items	Anthropic may use the Customer's name and logo publicly unless the Customer opts out via a form (clause G)		An NHS body should opt out of publicity on day one.	